

# Generative Adversarial Networks

Advanced Topics in Machine Learning

---

Andrea Cini, Daniele Zambon, Daniele Grattarola, Prof. Cesare Alippi

USI

Which photo is real?

---



# Which photo is real?



[thispersondoesnotexist.com](http://thispersondoesnotexist.com)

# Generative Adversarial Networks [1] in short

Goal: generate samples that **look like** the real data.

---

[1] I. J. Goodfellow et al., *Generative Adversarial Networks*, 2014.

# Generative Adversarial Networks [1] in short

Goal: generate samples that **look like** the real data.

Two main **neural networks**:

- **Generator  $G(z)$**  maps random noise  $z$  to the data space;
- **Discriminator  $D(x)$**  decides whether sample  $x$  was generated by  $G$  or is a real sample;

---

[1] I. J. Goodfellow et al., *Generative Adversarial Networks*, 2014.

# Generative Adversarial Networks [1] in short

Goal: generate samples that **look like** the real data.

Two main **neural networks**:

- **Generator**  $G(z)$  maps random noise  $z$  to the data space;
- **Discriminator**  $D(x)$  decides whether sample  $x$  was generated by  $G$  or is a real sample;

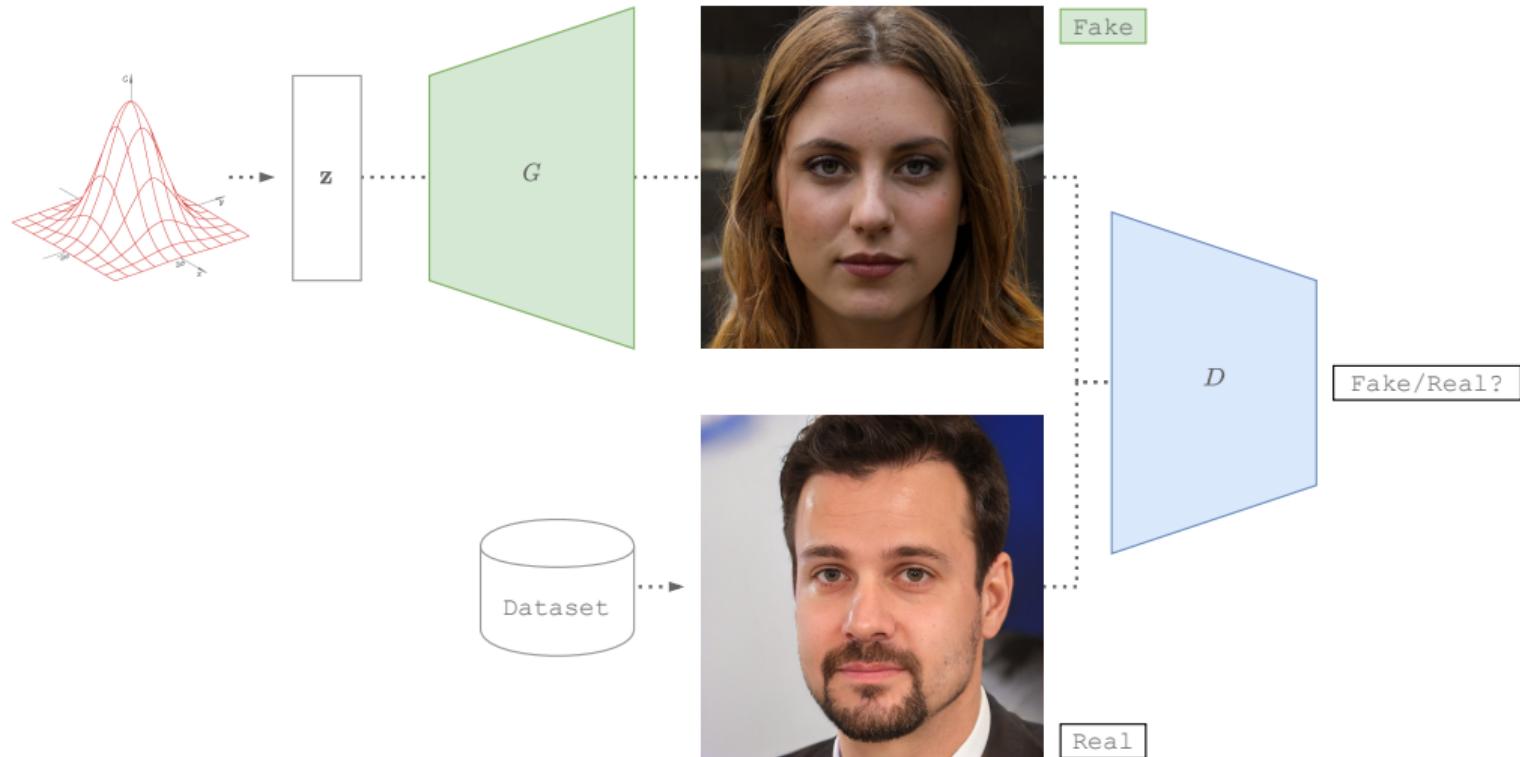
The two components play **against each other** until the **generator** fools the **discriminator**.

The diagram shows the text "The two components play **against each other** until the **generator** fools the **discriminator**." Below this text is a horizontal brace with a vertical line extending downwards from its center, labeled "Adversarial" below it. This visualizes the competitive relationship between the generator and discriminator.

---

[1] I. J. Goodfellow et al., *Generative Adversarial Networks*, 2014.

# Generative Adversarial Networks in short



Different types of neural networks can be used as **generator**:

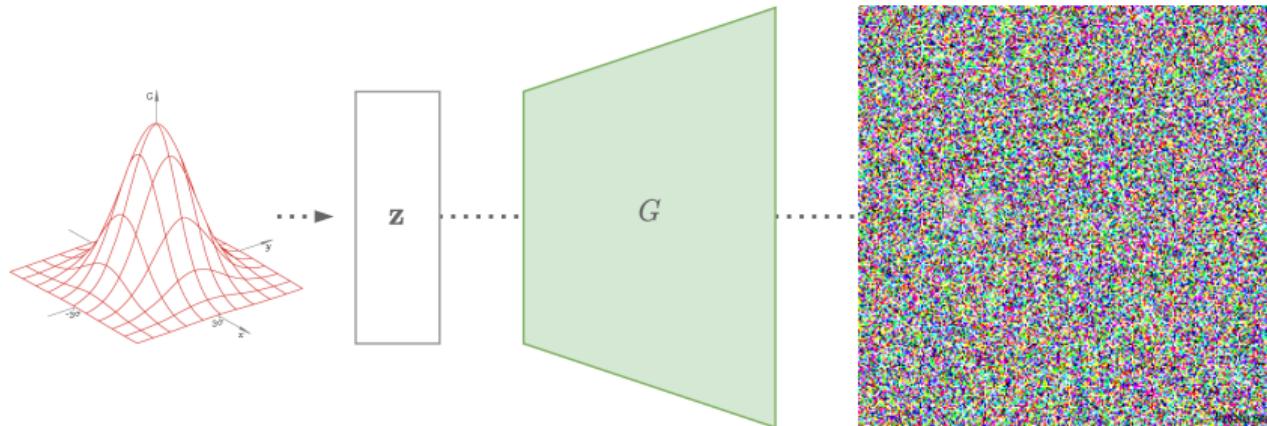
- Convolutional for images or audio;
- Recurrent for text or sequences;
- MLPs for tabular data;

Different types of neural networks can be used as **generator**:

- **Convolutional for images or audio;** ← we focus mostly on these (easier to visualize)
- Recurrent for text or sequences;
- MLPs for tabular data;

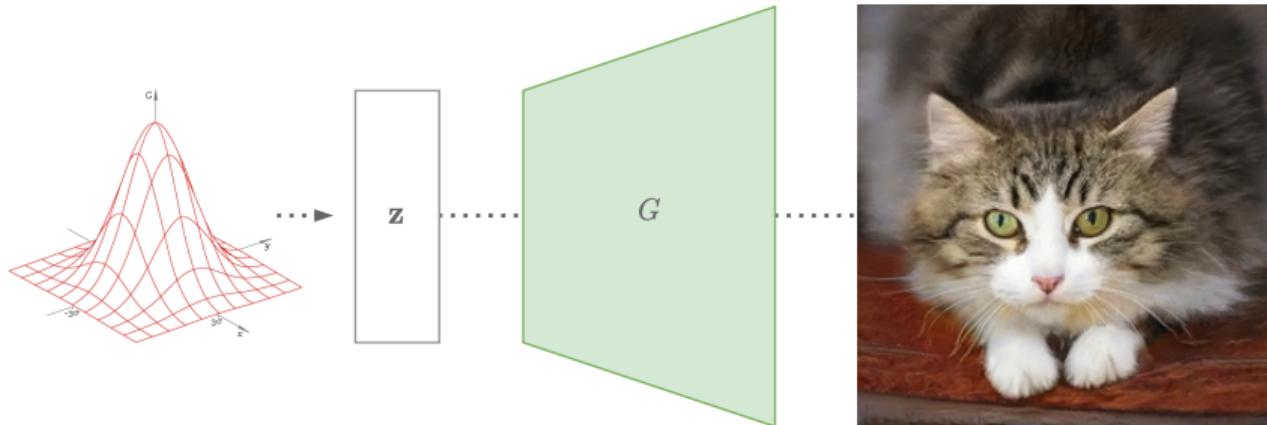
# Generator

Generates an image of  $w$  by  $h$  pixels  $\mathbf{x} = G(\mathbf{z}) \in [0, 1]^{w \times h}$ , where  $\mathbf{z} \sim p_z(\mathbf{z})$  (e.g.,  $\mathcal{N}(0, 1)$ ).



# Generator

Generates an image of  $w$  by  $h$  pixels  $\mathbf{x} = G(\mathbf{z}) \in [0, 1]^{w \times h}$ , where  $\mathbf{z} \sim p_z(\mathbf{z})$  (e.g.,  $\mathcal{N}(0, 1)$ ).



([thiscatdoesnotexist.com](http://thiscatdoesnotexist.com))

## Discriminator

---

The **discriminator** looks at samples from the data space  $\mathbf{x} \in [0, 1]^{w \times h}$  and outputs:

- 1 if the samples come from the real data distribution, *i.e.*,  $\mathbf{x} \sim p^*(\mathbf{x})$ ;
- 0 if the samples are **fake**, *i.e.*,  $\mathbf{x} \sim p_g(\mathbf{x})$  (first  $\mathbf{z} \sim p_z(\mathbf{z})$ , then  $\mathbf{x} = G(\mathbf{z})$ );

## Discriminator

---

The **discriminator** is trained to optimise two objectives:

- $\max_D \mathbb{E}_{x \sim p^*(x)} [\log D(x)]$  (output 1 on real samples)
- $\max_D \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]$  (output 0 on **generated** samples)

(this is just a different way to write a binary classification problem)

## Training the generator

Recall: the **generator** has to fool the **discriminator**.

$$\max_D \underbrace{\mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]}_{\text{Discriminator objective}}$$

## Training the generator

Recall: the **generator** has to fool the **discriminator**.

$$\underbrace{\max_D \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]}_{\text{Discriminator objective}} \rightarrow \underbrace{\min_G \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]}_{\text{Generator objective}}$$

Fooling = making the **discriminator** output 1 on generated samples.

## Putting everything together

---

If we combine the two objectives for  $G(\mathbf{z})$  and  $D(\mathbf{x})$  we get **the GAN min-max game**:

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p^*(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log (1 - D(G(\mathbf{z})))]$$

## Training the GAN

---

The GAN is trained by an iterative procedure, **repeated to convergence**:

1. Train the **discriminator** on a batch of real and fake samples;
2. Train the **generator** to fool the discriminator.

## Training the GAN - Discriminator training

For  $k$  steps do:

1. Sample a minibatch of noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from  $p_z(\mathbf{z})$ ;
2. Sample a minibatch of real samples from the dataset  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ ;
3. Update the weights  $\theta_D$  of  $D$  by gradient **ascent**:<sup>1</sup>

$$\nabla_{\theta_D} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log \left( (1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

---

<sup>1</sup>In practice we can also do gradient descent by changing the sign.

## Training the GAN - Generator training

---

Do once:

1. Sample a minibatch of noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from  $p_z(\mathbf{z})$ ;
2. Update the weights  $\theta_G$  of  $G$  by gradient **descent**:

$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m \log \left( (1 - D(G(\mathbf{z}^{(i)})) \right).$$

## Convergence

---

The min-max game goes on until  $G$  perfectly fools  $D$ .

# Convergence

The min-max game goes on until  $G$  perfectly fools  $D$ .

When this happens:

- The data distribution is **indistinguishable** from the distribution of the generated data:  
 $p^* = p_g$ .
- The global minimum of the training criterion is:

$$\mathbb{E}_{x \sim p^*(x)} \left[ \underbrace{\log D(x)}_{\log(0.5)} \right] + \mathbb{E}_{z \sim p_z(z)} \left[ \underbrace{\log (1 - D(G(z)))}_{\log(0.5)} \right] = \log(0.25)$$

# Convergence

The min-max game goes on until  $G$  perfectly fools  $D$ .

When this happens:

- The data distribution is **indistinguishable** from the distribution of the generated data:  
 $p^* = p_g$ .
- The global minimum of the training criterion is:

$$\mathbb{E}_{x \sim p^*(x)} \left[ \underbrace{\log D(x)}_{\log(0.5)} \right] + \mathbb{E}_{z \sim p_z(z)} \left[ \underbrace{\log (1 - D(G(z)))}_{\log(0.5)} \right] = \log(0.25)$$

i.e., the **discriminator** is **maximally confused** (can only output 0.5).

# Convergence

The min-max game goes on until  $G$  perfectly fools  $D$ .

When this happens:

- The data distribution is **indistinguishable** from the distribution of the generated data:  
 $p^* = p_g$ .
- The global minimum of the training criterion is:

$$\mathbb{E}_{x \sim p^*(x)} \left[ \underbrace{\log D(x)}_{\log(0.5)} \right] + \mathbb{E}_{z \sim p_z(z)} \left[ \underbrace{\log (1 - D(G(z)))}_{\log(0.5)} \right] = \log(0.25)$$

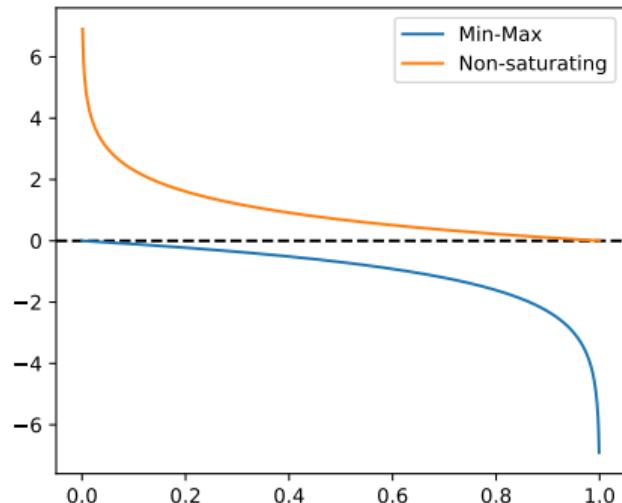
i.e., the **discriminator** is **maximally confused** (can only output 0.5).

**Note:** convergence is guaranteed given  $\underbrace{\text{sufficient capacity}}_{\text{Difficult to know...}}$  of the neural networks.

# Tips and Tricks

- Use **non-saturating** loss to optimize  $G$ :

$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(\mathbf{z}^{(i)}));$$

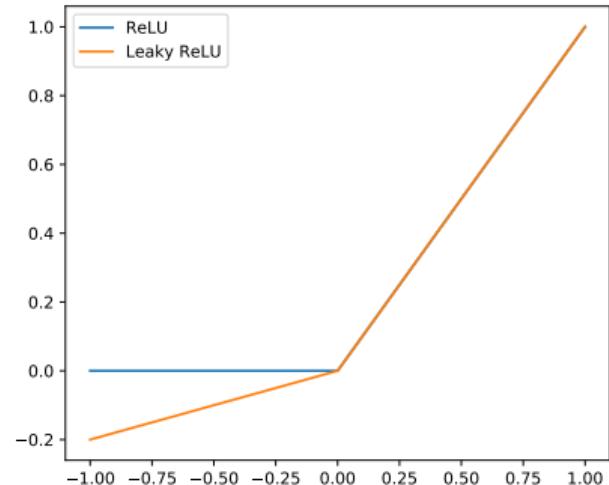


# Tips and Tricks

- Use **non-saturating** loss to optimize  $G$ :

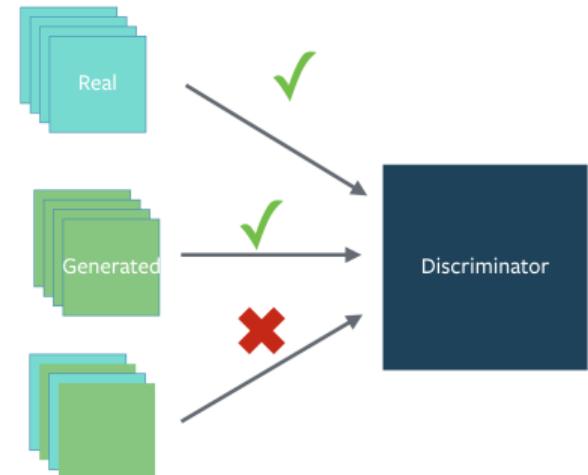
$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(\mathbf{z}^{(i)}));$$

- Avoid **sparse gradients**: use LeakyReLU, average pooling.



# Tips and Tricks

- Use **non-saturating** loss to optimize  $G$ :  
$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(\mathbf{z}^{(i)}));$$
- Avoid **sparse gradients**: use LeakyReLU, average pooling.
- **Do not mix** real and fake samples to train  $D$ ;



## Tips and Tricks

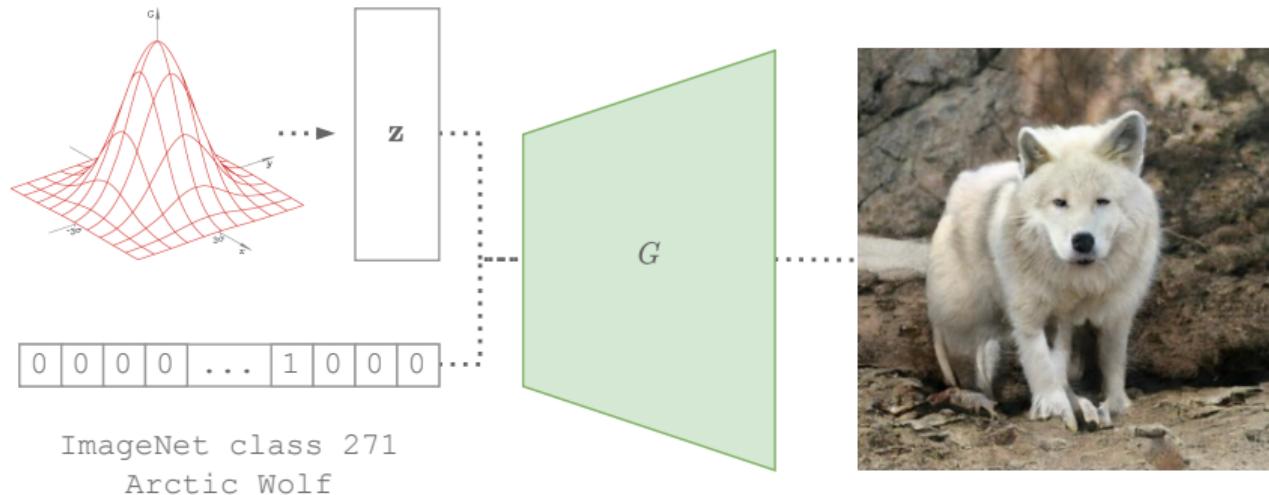
- Use **non-saturating** loss to optimize  $G$ :  
$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(\mathbf{z}^{(i)}));$$
- Avoid **sparse gradients**: use LeakyReLU, average pooling.
- **Do not mix** real and fake samples to train  $D$ ;
- Check out [github.com/soumith/ganhacks](https://github.com/soumith/ganhacks) for more ~~dark~~ magic empirical tips.



# Conditional GANs

In many cases, the samples from  $p^*$  are divided in classes (e.g., ImageNet).

Instead of generating **any** image from  $p^*$ , we generate  $\mathbf{x} = G(\mathbf{z}, y)$ , where  $y$  is a **class label**.



## Applications

---

## Text to image translation [2]

The small bird has a red head with feathers that fade from red to gray from head to tail

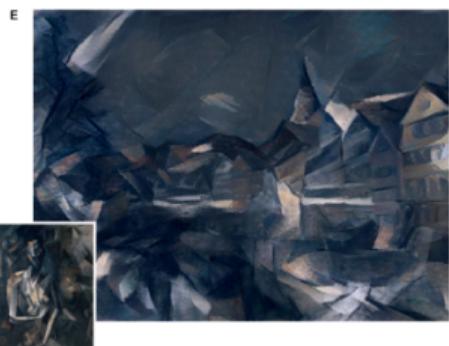


This bird is black with green and has a very short beak



[2] S. Reed et al., "Generative adversarial text to image synthesis," 2016.

# Style transfer [3]



[3] T. Karras et al., "A Style-Based Generator Architecture for Generative Adversarial Networks," 2018.

## Style transfer [4]



[4] G. Antipov et al., "Face aging with conditional generative adversarial networks," 2017.

## Super-resolution [5]



---

[5] C. Ledig et al., "Photo-realistic single image super-resolution using a generative adversarial network," 2016.

## Demo

---

## Demo

---



[cutt.ly/sfO5CU9](https://cutt.ly/sfO5CU9)

Questions?

## References i

- [1] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, *Generative adversarial networks*, 2014. arXiv: 1406.2661 [stat.ML].
- [2] S. Reed, Z. Akata, X. Yan, L. Logeswaran, B. Schiele, and H. Lee, “Generative adversarial text to image synthesis,” *arXiv preprint arXiv:1605.05396*, 2016.
- [3] T. Karras, S. Laine, and T. Aila, “A style-based generator architecture for generative adversarial networks,” in *CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 4396–4405.
- [4] G. Antipov, M. Baccouche, and J.-L. Dugelay, “Face aging with conditional generative adversarial networks,” in *2017 IEEE international conference on image processing (ICIP)*, IEEE, 2017, pp. 2089–2093.

- [5] C. Ledig, L. Theis, F. Huszár, J. Caballero, A. Cunningham, A. Acosta, A. Aitken, A. Tejani, J. Totz, and Z. Wang, "Photo-realistic single image super-resolution using a generative adversarial network," *arXiv preprint arXiv:1609.04802*, 2016.